

Review questions

Edward L. Ionides, Aaron A. King, Qianying Lin

June 27, 2021

Questions on Lesson 2: Simulation of stochastic dynamic models

Question 2.1. Let V_n be a Markov process and let $W_n = h(V_n)$ for some function h . Which of the following statements are true?

- i) W_n is a Markov process for all choices of h .
- ii) W_n is a Markov process for some choices of h .
- iii) W_n is not a Markov process for any choice of h .
- iv) If $V_n = (X_n, Y_n)$ where X_n and Y_n are a POMP model, and $h(X_n, Y_n) = X_n$ then W_n is a Markov process.
- v) If $V_n = (X_n, Y_n)$ where X_n and Y_n are a POMP model, and $h(X_n, Y_n) = Y_n$ then W_n is a Markov process.

- (A) i,iv,v
- (B) ii,iv
- (C) ii,v
- (D) iii
- (E) None of the above

Solution. (B)

Questions on Chapter 3: Likelihood for POMP

Question 3.1. Suppose that 10 replications of a particle filter, each using 10^3 particles, runs in 15 minutes with no parallelization. To look for a more precise likelihood evaluation, you consider running 20 replicates, each with 10^4 particles. How many minutes will this take, if you distribute the calculation across 4 cores?

- (A) 50
- (B) 60
- (C) 75
- (D) 120
- (E) 300

Solution. (C). Using the linear dependence, also called proportionality, of the computing effort on various algorithmic parameters, we calculate

$$5 \times (10000/1000) \times (20/10) \times (1/4) = 75.$$

Question 3.2. A particle filter is repeated 5 times to evaluate the likelihood at a proposed maximum likelihood estimate, each time with 10^4 particles. Suppose the log likelihood estimates are -2446.0 , -2444.0 , -2443.0 , -2442.0 , -2440.0 . Which of the following is an appropriate estimate for the log likelihood at this parameter value and its standard error.

- (A) Estimate = -2443.0 , with standard error 1.0
- (B) Estimate = -2443.0 , with standard error 2.2
- (C) Estimate = -2443.0 , with standard error 5.0
- (D) Estimate = -2441.4 , with standard error 2.2
- (E) Estimate = -2441.4 , with standard error 1.0

Solution. (E). Answers (A), (B) and (C) estimate using a mean on the log scale. However, the particle filter provides an unbiased likelihood estimate on a natural scale but not on a log scale. Note that the particle filter also has some bias for most quantities on a natural scale, which reduces to zero as the number of particles tends to infinity, but it happens to be unbiased for the likelihood. The standard error for the log of the mean of the likelihoods can be computed by the delta method or a jack-knife, for example using the `logmeanexp` function in `pomp`.

```
l1 <- c(-2446, -2444, -2443, -2442, -2440)
mean(l1)

[1] -2443

sd(l1)

[1] 2.236068

sd(l1)/sqrt(length(l1))

[1] 1

library(pomp)
logmeanexp(l1, se=TRUE)

               se
-2441.422198    1.380688
```

Question 3.3. What is the log likelihood (to the nearest unit) of the Dacca cholera data for the POMP model constructed in `pomp` via

```
d <- dacca(deltaI=0.08)
```

with cholera mortality rate 8% and other parameters fixed at the default values.

- (A) -3764
- (B) -3765
- (C) -3766
- (D) -3767
- (E) -3768

Solution. (A), calculated as follows:

```
d <- dacca(deltaI=0.08)
library(doParallel)
my_cores <- detectCores()
registerDoParallel(my_cores)
bake(file="Q9.rds",{
  foreach(i=1:32,.combine=c) %dopar% {
    library(pomp)
    logLik(pfilter(d,Np=10000))
  }
}) -> cholera_loglik
logmeanexp(cholera_loglik,se=TRUE)

                                se
-3764.2030946      0.1502043
```

Questions on Chapter 4: Inference via iterated filtering

Question 4.1. When carrying out inference by iterated particle filtering, the likelihood increases for the first 10 iterations or so, and then steadily decreases. Testing the inference procedure on simulated data, this does not happen and the likelihood increases steadily toward convergence. Which of the following is the best explanation for this?

- (A) One or more random walk standard deviation is too large.
- (B) One or more random walk standard deviations is too small.
- (C) The model is misspecified, so it does not fit the data adequately.
- (D) A combination of the parameters is weakly identified, leading to a ridge in the likelihood surface.
- (E) Too few particles are being used.

Solution. (C). All the other issues can cause inference problems, but likely would cause similar problems on simulated data.

When there is a reproducible and stable phenomenon of decreasing likelihood, it generally indicates that the unperturbed model is a worse fit to the data than the perturbed model. Recall that the likelihood calculated by iterated filtering at each iteration corresponds to the model with perturbed parameters rather than the actual postulated model with fixed parameters. If the perturbed model has higher likelihood, it may mean that the data are asking to have time-varying parameters. It may also be a signature of any other weakness in the model that can be somewhat accommodated by perturbing the parameters.

Question 4.2. People sometimes confuse likelihood profiles with likelihood slices. Suppose you read a figure which claims to construct a profile confidence interval for a parameter ρ in a POMP model with four unknown parameters. Which of the following confirms that the plot is, or is not, a properly constructed profile confidence interval. The code producing the plot is available to you as an Rmarkdown file.

- (A) The CI is constructed by obtaining the interval of ρ values whose log likelihood is within 1.92 of the maximum on a smoothed curve of likelihood values plotted against ρ .
- (B) The code involves evaluation of the likelihood but not maximization.
- (C) The points along the ρ axis are not equally spaced.
- (D) The smoothed line shown in the plot is close to quadratic.
- (E) (A) and (D) together.

Solution. (B).

If the researchers calculate a sliced likelihood through the MLE and tell you it is a profile, but you are concerned they might have constructed a slice by mistake, it is hard to know without looking at the code. (A) is the proper construction of a profile if the points are maximizations over the remaining parameters for a range of fixed values of ρ . However, if the code does not involve maximization over other parameters at each value of ρ , it cannot be a proper profile. It could be a slice accidentally explained to be a profile, and with a confidence interval constructed as if it were a profile.

Questions on Chapter 5: Measles

Question 5.1. Two models are fitted to case counts on an epidemic. Model 1 is an SIR POMP model with a negative binomial measurement model, and model 2 is a linear regression model estimating a cubic trend. The log likelihoods are $\ell_1 = -2037.91$ and $\ell_2 = -2031.28$ respectively. Which of the following do you agree with most?

- (A) We should not compare the models using these likelihoods. They correspond to different model structures, so it is an apples-to-oranges comparison.
- (B) We can compare them, but the difference is in the 4th significant figure, so the likelihoods are statistically indistinguishable.
- (C) The linear model has a noticeably higher likelihood. Our mechanistic model needs to be updated to beat this benchmark before we can responsibly interpret the fitted model. If a simple regression model has higher likelihood than a more complex mechanistic model, one should prefer the simpler model.
- (D) The linear model has a noticeably higher likelihood. The mechanistic model is somewhat validated by being not too far behind the simple regression model. We are justified in cautiously interpreting the mechanistic model, while continuing to look for further improvements.
- (E) The log likelihoods cannot properly be compared as presented, but could be if we used a Gaussian measurement model for the POMP (or a negative binomial generalized linear model instead of least squares for the regression).

Solution. (D).

Why not (A)? Likelihoods of different models for the same data can be compared. Likelihood ratio tests using Wilks's theorem specifically require nested models, but in other contexts (such as AIC and the Neyman-Pearson lemma) the models being compared by likelihood do not need to have any particular relationship. Why not (B)? Likelihood ratios have statistical meaning, which corresponds to differences of log likelihoods. The likelihood is a dimensional quantity, whereas the likelihood ratio is dimensionless. The units used correspond to a scientifically arbitrary additive constant to the log likelihood, which disappears after taking differences.

Why not (C)? If our only goal were to find a predictive model, then (C) could be a reasonable position. Usually, we want to find a model that also has interpretable structure, leading to understanding of the system or estimating the effect of interventions. A simple regression model cannot do those things, even if it fits a bit better. If the mechanistic model fits much worse than simple alternatives, it is not providing a reasonable explanation of the data, suggesting that there may be important things missing from the model specification. Quite likely, with some persistence, a mechanistic specification will beat a simple off-the-shelf statistical model.

Question 5.2. A compartment model is first implemented as a system of ordinary differential equations (ODEs). This leads to qualitatively reasonable trajectories, but poor likelihood values. The researchers add stochasticity in an attempt to improve the fit of the model by interpreting the ODEs as rates of a Markov chain. The likelihood, maximized by iterated particle filtering, remains poor compared to ARMA benchmarks. In addition, the effective sample size for the particle filtering is low at many time points despite even using as many as 10^4 particles. Which of the following is the most promising next step.

- (A) Increase to 10^5 particles, moving the computations to a cluster if necessary.
- (B) Add noise to one or more rates to allow for overdispersion.
- (C) Try adding extra features to the model to capture scientific details not present in the original model.
- (D) Experiment with variations in the iterated filtering procedure; maybe more iterations, or a different cooling schedule.
- (E) To address the possibility of reporting errors, see if the model fits better when the most problematic data points are removed.

Solution. (B). All the possibilities are worth consideration. However, adding noise in rates to give flexibility in mean-variance relationships is commonly an important part of developing a stochastic model. The simple compartment model interpretation of a ODE as a Markov chain is determined by the rates and therefore does not have free parameters to describe variance. There is some variance inherent in the Markov chain (demographic stochasticity) but additional variability may be needed. It will be hard to investigate the other possibilities if the model has not been given enough stochasticity to explain the variability in the data, so including overdispersion should be an early step. Note that overdispersion can be included in both the process model and the measurement model.

Question 5.3. You fit an SEIR model to the case report data on an immunizing disease for a city. The resulting confidence interval for the mean latent period is 12-21 days, but clinical evidence points to a latent period averaging about 7 days. Which of the following is the most appropriate response to this discrepancy.

- (A) The latent period may be confounded with some unmodeled aspect of the system, such as spatial or age structure. The model estimates an effective latent period at the population level, which may not perfectly match what is happening at the scale of individuals. One should be cautious of making a causal interpretation of models fitted to observational data.
- (B) The discrepancy shows that something is substantially wrong with the model. Extra biological detail should be introduced with the goal of bringing the estimated parameter back in line with the known biology of the system.
- (C) The discrepancy is problematic, but fortunately can readily be fixed. Since we know the clinical value of this parameter with reasonable accuracy, we should simply use this value in the model rather than estimating it.
- (D) If the model fits the data statistically better than any known alternative model, then we have to take the estimated parameter at face value. It is certainly possible that the estimates in the literature correspond to some different population, or different strain, or have some other measurement bias such as corresponding to severe cases resulting in hospitalization. The discrepancy does not show that our model was wrong.
- (E) This discrepancy suggests that we should take advantage of both C and D above by putting a Bayesian prior on the latent period. By quantifying the degree of our skepticism about the previously established clinical value of 7 days, we can optimally combine that uncertainty with the evidence from this dataset.

Solution. (A).

Transferring parameter estimates between scales is hard. An example is the difficulty of reconciling micro and macro economics. It is generally not possible to guarantee that a parameter means exactly the same thing in models at different scales. (A) acknowledges this. The other answers, in various ways, assume that there should be a single parameter value that describes the system at all scales.